

Naive Bayes

☺ Type	Lecture Notes
☑ Reviewed	<input type="checkbox"/>
⚙ Available Summary?	Done
# Week	5

YOU SHOULD KNOW:

- training and using classifier based on Bayes rule
- conditional independence:
 - what it is
 - why is it important
- Naive Bayes:
 - what it is
 - how to estimate the parameters
 - how to make predictions

▼ Bayes Optimal Classifier - Part 1

- Goal is to learn $f : X \rightarrow Y$
 - $X \rightarrow$ features
 - $Y \rightarrow$ target class

$$Y^* = \operatorname{argmax}_{y_k} P(Y = y_k | X)$$

▼ how many parameters must we estimate?

▼ boolean and 2 features

Suppose $X = (X_1, X_2)$ where X_i and Y are boolean RV's, to estimate $P(Y | X_1, X_2)$:

X_1	X_2	$P(Y = 1 X_1, X_2)$	$P(Y = 0 X_1, X_2)$
0	0	0.1	0.9
1	0	0.24	0.76
0	1	0.54	0.46
1	1	0.23	0.77

4 parameters!

$$P(Y = 0 | X) = 1 - P(Y = 1 | X)$$

▼ boolean and n features

Suppose $X = (X_1, X_2, \dots, X_n)$ where X_i and Y are boolean RV's, to estimate $P(Y | X_1, X_2, \dots, X_n)$:

we need 2^n rows to create the table

▼ can we reduce parameters using Bayes Rule?

Suppose $X = [X_1, \dots, X_n]$ and all X_i s and Y are boolean variables.

To estimate $P(X_1, \dots, X_n | Y)$:

2^n rows to create the table and -1 cause all the probabilities are sum up to 1 (so we already know the final probability we do not need to calculate)

$2 * (2^n - 1)$ parameters we need

To define $P(Y)$:

we need 1 parameter

Finally:

we need $2 * (2^n - 1) + 1$ parameters

▼ Naive Bayes

- Naive bayes assumes:

$$P(X_1, \dots, X_n|Y) = \prod_i P(X_i|Y)$$

that X_i and X_j are conditionally independent given Y , for all $i \neq j$

- Naive bayes uses assumption that the X_i are conditionally independent given Y

$$P(X_1, X_2|Y) = P(X_1|X_2, Y).P(X_2|Y)$$

more generally;

$$\underbrace{P(X_1, \dots, X_d|Y)}_{\text{likelihood}} = \prod_{i=1}^d P(X_i|Y)$$



Recall conditionally independence:

X is conditionally independent of Y given Z , if the probability distribution governing X is independent of the value Y , given the value of Z

$$P(X = i|Y = j, Z = k) = P(X = i|Z = k)$$

$$P(X, Y|Z) = P(X|Z).P(Y|Z)$$

▼ How many parameters to describe $P(X_1, \dots, X_n|Y)$ and $P(Y)$?

Without conditional indep assumption $\rightarrow 2 * (2^n - 1)$ and 1

With conditional indep assumption $\rightarrow 2n$ and 1

▼ Naive Bayes classification

Bayes Rule:

$$P(Y = y_k|X_1, \dots, X_n) = \frac{P(Y=y_k).P(X_1, \dots, X_n|Y=y_k)}{\sum_j P(Y=y_j).P(X_1, \dots, X_n|Y=y_j)}$$

Assuming conditional independence among X_i 's:

$$P(Y = y_k|X_1, \dots, X_n) = \frac{P(Y=y_k). \prod P(X_i|Y=y_k)}{\sum_j P(Y=y_j). \prod P(X_i|Y=y_j)} \quad \text{estimate in training}$$

So, to pick most probable Y for $X^{new} = (X_1, \dots, X_n)$:

$$Y^{new} \leftarrow \underset{y_k}{\operatorname{argmax}} P(Y = y_k) \prod_i P(X_i^{new}|Y = y_k) \quad \text{testing}$$

▼ EXAMPLE

- Train Naïve Bayes (examples)

for each* value y_k

estimate $\pi_k \equiv P(Y = y_k)$

for each* value x_{ij} of each attribute X_i

estimate $\theta_{ijk} \equiv P(X_i = x_{ij}|Y = y_k)$

- Classify (X^{new})

$$Y^{new} \leftarrow \underset{y_k}{\operatorname{argmax}} P(Y = y_k) \prod_i P(X_i^{new}|Y = y_k)$$

$$Y^{new} \leftarrow \underset{y_k}{\operatorname{argmax}} \pi_k \prod_i \theta_{ijk}$$



Posterior = Likelihood * Prior

In naive bayes, we are multiplying lots of small numbers. It may cause underflow.



Underflow occurs when you perform an operation that's smaller than the smallest magnitude non-zero number.

Solution:

Perform all computations by summing logs of probabilities rather than multiplying probabilities.

$$p_1 * p_2 = e^{\log(p_1) + \log(p_2)}$$

▼ Document Classification using Naive Bayes:



To generate a document for giving class; randomly pick a number to predict for, if the $word_i$ should be in the document or not.

▼ Bernoulli document model:

a document is represented by a **binary feature vector**, whose elements indicate **absence or presence of corresponding word** in the document

▼ Classification with Bernoulli Model

Training data:

Matrix: X

document: i

feature vector: X_i

presence of word t in the document i : X_{it}

Parameter Estimation (MLE):

Priors: $P(Y = y_k) = \frac{N_k}{N}$

Likelihoods: $P(w_t|Y = y_k) = \frac{n_k(w_t)}{N_k}$

(use a likelihood vector)

Classify new document D with feature vector X :

$$P(Y = y_k|X) = P(Y = y_k) \cdot P(X|Y = y_k)$$

$$P(X|Y = y_k) = \prod_{t=1}^{|V|} [X_t \cdot P(w_t|Y = y_k) + (1 - X_t)(1 - P(w_t|Y = y_k))]$$

▼ Multinomial document model:

a document is represented by an **integer feature vector**, whose elements indicate **frequency of corresponding word** in the document

▼ Classification with Multinomial Model

Training data:

Matrix: X

document: i

feature vector: X_i

the count of number of times w_t occurs in document i : X_{it}

1 if document i is of class y_k , otherwise 0: $Z_{i,k}$

Parameter Estimation (MLE):

Priors: $P(Y = y_k) = \frac{N_k}{N}$

Likelihoods: $P(w_t|Y = y_k) = \frac{\sum_{i=1}^N X_{i,t} \cdot Z_{i,k}}{\sum_{t'=1}^{|V|} \sum_{i=1}^N X_{i,t'} \cdot Z_{i,k}}$

The relative frequency of w_t in documents of class $Y = y_k$ with respect to the total number of words in documents of that class

Classify new document D with feature vector X :

$$P(Y = y_k|X) = P(Y = y_k) \cdot P(X|Y = y_k)$$

$$P(X|Y = y_k) = \prod_{t=1}^{|V|} P(w_t|Y = y_k)^{X_t}$$

Parameter estimation (MAP, with $\alpha = 1$):

Likelihoods: $P(w_t|Y = y_k) = \frac{\alpha + \sum_{i=1}^N X_{i,t} \cdot Z_{i,k}}{|V| \alpha + \sum_{t'=1}^{|V|} \sum_{i=1}^N X_{i,t'} \cdot Z_{i,k}}$

A trick to avoid zero counts which is equivalent to MAP estimation with Dirichlet prior with $\alpha = 1$.

▼ Bayes Optimal Classifier - Part 2

- Very fast, low storage requirements
- robust to irrelevant features

- optimal if the conditional independence assumptions hold
- a good dependable baseline for text classification

▼ **Violating the NB classifier:**

- usually features are not conditionally independent
- in practice it often works well
- it does not produce accurate probability estimates when its independence assumptions are violated, it may still (and often) pick the correct maximum-probability class in many cases
- typically handles noise well since it does not even focus on completely fitting the training data